

Assignment 2:Text Classification

September 2020

1 Introduction

Stack Exchange is a very popular Q&A (Question-and-Answer) based website. We want to analyze some archived data of Stack Exchange using text classification. The link to stack exchange archive is

<https://archive.org/details/stackexchange>

The goal of text classification is to identify the topic for a piece of text (news article, web-blog, etc.). Text classification has obvious utility in the age of information overload, and it has become very popular for applied machine learning algorithms. In this project, you will implement k-nearest neighbor and Naive Bayes, apply these to text classification on Stack Exchange sample data, and compare the performances of these techniques.

2 Description of the Dataset

For your convenience, we have collected and accumulated a small replica of stack exchange dataset, on which you will perform our analysis. The description of the dataset is give below:

1. The size of the dataset is more 100MB. Download the zipped file containing data. The link of the dataset that we shall use is :

<https://www.dropbox.com/s/1jdct708qk8p6za/Data.zip?dl=0>

2. Just consider the training files except (3d_printer.xml). Although the name of the folder is training, the same files under this folder will be used for train, validation and test purpose. For each of the topic (as found in topic.txt), take the first 500 texts/rows/documents (herein we shall refer a line/row as document) of the file for training purpose, next 200 documents for validation purpose and finally next 500 documents for test purpose. So, if you have 3 topics, then your training size 1500 documents, validation size will be 600 documents and test size will be 1500 documents.

3. For training, validation and test purpose, take every line which starts with "row" and keep only the "Body" portion of this row. **Consider only this portion as a document and the name of the file as the topic name.**

3 Text Preprocessing

The extracted text should be preprocessed before actual use. Following operations, at least, have to be performed as part of text preprocessing.

1. Conversion to lowercase
2. Punctuation removal
3. Stopword removal
4. Tokenization
5. Stemming
6. Lemmatization

4 k-Nearest Neighbor (k-NN)

1. Implement the k-NN algorithm for text classification. Try the following distance or similarity measures with their corresponding representations:
 - Hamming distance: each document is represented as a boolean vector, where each bit represents whether the corresponding word appears in the document.
 - Euclidean distance: each document is represented as a numeric vector, where each number represents how many times the corresponding word appears in the document (it could be zero).
 - Cosine similarity with TF-IDF weights (a popular metric in information retrieval): each document is represented by a numeric vector as in the case of euclidean distance. However, now each number is the TF-IDF weight for the corresponding word (as defined below). The similarity between two documents is the dot product of their corresponding vectors, divided by the product of their norms.
2. Let w be a word, d be a document, and $N(d, w)$ be the number of occurrences of w in d (i.e., the number in the vector as in the case of euclidean distance). TF stands for term frequency, and $TF(d, w) = N(d, w)/W(d)$, where $W(d)$ is the total number of words in d . During TF calculation, You can simply omit the word when the word is new in test/validation set. IDF stands for inverted document frequency, and $IDF(d, w) = \log(\frac{D+\alpha}{C(w)+\beta})$, where D is the total number of documents, $C(w)$ is the total number of documents that contains the word w , α and β are arbitrary values to

avoid the IDF from zero value and division-by-zero respectively (During implementation, you can set $\alpha = \beta = 0$, and put a very small constant when the word is available in all training documents. On the other hand, you can simply omit the word when the word is new in test/validation set). The base for the logarithm is not the determining factor, you can use e or 2. The TF-IDF weight for w in d is $TF(d, w) * IDF(d, w)$; this is the number you should put in the vector in Cosine similarity. TF-IDF is a clever heuristic to take into account of the "information content" that each word conveys, so that frequent words like "the" is discounted and document-specific ones are amplified. You can find more details about it online or in standard IR text.

3. You should try $k = 1$, $k = 3$ and $k = 5$ with each of the representations above. Notice that with a distance measure, the k -nearest neighborhoods are the ones with the smallest distance from the test point (of validation set or test set), whereas with a similarity measure, they are the ones with the highest similarity scores.
4. Output the result of running all the three values of k using all the three k -NN techniques into a single file.

5 Naive Bayes

Implement the Naive Bayes algorithm for text classification. Naive Bayes used to be the de facto method for text classification.

1. Consider all the words of a test (or validation) document independently, then calculate the probability of the document of being a topic, and then pick up the topic which provides the highest probability score.
2. Try different smoothing factors (at least 10 different values) and calculate the accuracy for each value of smoothing factor.

6 Comparison and Report Writing

In this part, you will compare between the performance of k -NN classifier and Naive Bayes classifier for text classification. Follow the steps below:

1. Write down the output on validation set (A matrix containing methodologies along rows and three different k values along columns for k -NN, and a matrix containing 10 different values of smoothing factor along the rows for NB. Each cell will contain the corresponding accuracy value.)
2. Take the best classifier (considering both the measure and value of k) from k -NN. Also, take the Naive Bayes set with the best smoothing factor value. Run 50 times both the K -NN and Bayesian learner on the test set (for each iteration, take 10 documents from each of the topics). Write

down the accuracy values for each of the 50 iterations on test set (both for k-NN and NB).

3. Compute t-statistic at significance levels of 0.005, 0.01 and 0.05, and compare which algorithm (k-NN or Bayesian) is better. Write the results in a report as well as the justification of the result in your own words.
4. Search in the Internet to learn t-statistics in details.
5. Never copy the report. Just answer the questions precisely. Make it as simple as possible. Too much description is not needed!

7 Assignment Summary

For each of the above-mentioned algorithms (k-NN and NB), just consider the training files except (3d_printer.xml). Although the name of the folder is training, the same files under this folder will be used for train, validation and test purpose. For each of the topic (as found in topic.txt), take the first 500 documents of the file for training purpose, next 200 documents for validation purpose and finally next 500 documents for test purpose. So, if you have 3 topics, then your training size 1500 documents, validation size will be 600 documents and test size will be 1500 documents. Tune the hyperparameters and select the available measures. For example, the best combinations of the distance measure (Hamming distance, Euclidean distance and Cosine similarity with TF-IDF) and the values of hyperparameter k (k=1,3 and 5) in case of k-NN. On the other hand, select the smoothing factor value that provides the best result in case of NB. Use the validation set to select the best combination for k-NN and best smoothing factor value for NB. After that, do the test and compare. For k-NN, take the best distance measure and k value combination, then perform 50 iterations on test set and generate the accuracy in each iteration. In an iteration, take 10 documents from each topic (i.e., 30 if the topic number is 3), count the performance (accuracy) for this iteration and save it in a file. Store the accuracy values of 50 iterations in the same file. Repeat this procedure for NB where you will generate another file containing 50 accuracy values for 50 iterations. The contents of these two values will be your input to t-statistics. Perform the t-statistics, generate the result and write down the result in the report. So, basically your report will contain 1) the output on validation set (A matrix containing methodologies along rows and three different k values along columns for k-NN, and a matrix containing 10 different values of smoothing factor along the rows for NB. Each cell will contain the corresponding accuracy value), 2) the accuracy values of 50 iterations on test set (both for k-NN and NB) and 3) the calculation and result of t-statistics along with your justification.

8 Special Instructions

1. You can use built-in libraries for file reading, text preprocessing and t-statistics analysis. But you CANNOT use any dedicated library for k-NN and Naive Bayes implementation. You have to implement these from scratch. However, you can use built-in utility libraries (e.g., numpy in python) during k-NN and Naive Bayes implementation.
2. Don't Copy anything! If you do copy from internet or from any other person or from any other source, you will be severely punished. More than that, we expect Fairness and honesty from you. Don't disappoint us!
3. The report should be in .docx/.pdf. Write precisely in your own language and keep it as simple as possible.
4. Your submission must have two files- one file containing the code and another one .docx/.pdf file for report. Every file must have your student id as prefix. For example, "1505125.Code.py" and "1505125.Report.pdf". Note that you don't need to upload the data file.
5. The deadline of the submission is 11:55 PM, October 12, 2020 (Monday). Considering the technical problems, you should not submit the assignment at the last moment.

9 About Version

A version is included with the assignment pdf name. This is because if any information is changed, then the version will be upgraded and the changes will be summarized in this section.

9.1 Changes included in Version 2

1. A new section (Section 7) named "Assignment Summary" has been added. You should go through this section first.
2. "Description of the Dataset" section has been updated.
3. "Nearest Neighbor (k-NN)" section has been updated.
4. "Naive Bayes" section has been updated.
5. "Comparison and Report Writing" section has been updated.
6. Only the 4th point of the "Special Instructions" section has been updated.